


H20 / B300 runtime: establish a faithful baseline before kernel work

ILLUSTRATIVE · NOT MEASURED


CURRENT EXPERIMENT PATH

Logical execution flow · not a GPU topology or timeline



FUTURE WORK · NOT IMPLEMENTED

Proceed only if profiling identifies a material bottleneck



No H20/B300 speedup is established. A sparse mask alone does not prove skipped computation or lower latency.