


DeepProbe-VLM: one visual pass, a genuinely recurrent language model

ILLUSTRATIVE · NOT MEASURED



Inference parameters stay fixed.

Visual encoding once $\neq$ skipping visual-token updates inside the LM.

Current baseline: dense native loops. Selective refresh and custom CUDA kernels are future work, not shown as implemented.